

UBT v2.1

Universal Basic Time Workload, capacity, and recovery

CC0 1.0 Universal - Public Domain. No attribution required.

Status: Reference specification with Geometry v2.6 integration. Operational profiles are proposals; source mathematical results retain their declared domains. No implementation or field-test result is reported here.

UBT makes the work of keeping a commons running visible: creation, care, repair, coordination, training, and deferred maintenance. Its purpose is to reduce avoidable burden and protect the capacity to continue.

A time record describes work. It does not measure human worth or grant political standing, governance weight, food, water, rank, or permanent entitlement. Essential support does not depend on contributing or agreeing to personal time tracking.

The four retained stewardship practices

Recognition Pulse. Offer prompt, consent-compatible acknowledgment. Measure delay if useful. Acknowledgment does not certify the work or require praise, emotional disclosure, or a public identity.

Sabbath of the Commons. Protect recovery through realistic scheduling, coverage, and the ability to decline more work. The obligation falls on the work arrangement to respect agreed limits; it does not authorize control of a person's private time.

Mirroring Scribe. Reflect back what was done or requested so errors can be corrected with little ceremony. Keep acknowledgment distinct from verification of a technical claim.

Pancake Report. Keep a short, usable field note: what happened, what worked, what failed, and what would reduce future effort. Preserve a source or version link where it matters.

The older promise that every unit of presence meets a unit of peace remains an aspiration. Whether people obtain more usable time must be observed, not inferred from the ritual names.

Work records and capacity

The proposed UBT-Work/2.1 profile can be a simple local record. Use only the detail needed for the stated coordination purpose. A person may decline personal logging; missing records do not mean that no work or care occurred.

Field group	What to retain
Identity and scope	Stable task ID, source namespace, record version, purpose, system or asset reference, and observation window.
Work and time	Task type; estimate or observation; units, method and uncertainty; actual/deferred state; output, upkeep and capacity references.
Capacity and coverage	Required skill, voluntary availability, agreed limits, supervision, current coverage, backup, and training or replacement time.
Evidence and correction	Source, OSIIN version and UPRR references; Geometry-Claim/1 record; completion criterion; correction and next affected decision.
Privacy and custody	Consented identifier or no identifier; public aggregate fields; restricted fields; local custodian; retention, export, and withdrawal rules.

Count the work that supports the visible output

Include preparation, documentation, validation, cleaning, transport, administration, training, supervision, maintenance, repair, and care when they fall within the declared task scope. Teaching, childcare, eldercare, mediation, meals, and relational work can be represented without exposing intimate details.

Keep different kinds of time separate

Person-hours and elapsed time answer different questions. Two people working one hour produce two person-hours and one elapsed hour. Keep estimates, observations, volunteered capacity, and recovery opportunity in separate fields. Unknown is not zero.

Make shortages visible without assigning blame

Review backlog age, unfilled work, overdue inspections, deferred leave, and reliance on one skill holder. Compare maintenance load and voluntary capacity only within compatible tasks, skills, conditions, and intervals. A surplus of unrelated hours does not replace a missing specialist.

Knowledge and resources in practice

The eight companion practices from the original tuning document remain available in revised form. They support the separate functions defined by OSIIN v3.1 and UPRR v3.1.

UPRR practices

Seamless Unlock. Reduce avoidable request and delivery delay. Visibility may prompt a request; actual allocation follows the declared local or bridge rule. Recognition never becomes an automatic condition for essentials.

Use-as-Offering. Offer clear care and return instructions. Record use, wear, repair, and replenishment separately; using a resource does not automatically improve it or create a work credit.

Regeneration Feedback. Send relevant maintenance and waste observations to OSIIN. Preserve the measurement basis and uncertainty. A proposed improvement is not yet observed regeneration.

Sovereignty Shield. Keep disclosure voluntary, custody explicit, rules inspectable, and alternatives practical. Export or fork the records and procedures where permitted; a file fork does not duplicate a physical resource.

OSIIN practices

Sanitation of the Field. Correct obsolete guidance and simplify duplicate views while preserving distinct failures, source material, and version history. No fixed deletion quota applies.

Propagation of the Lineage. Use documentation, teaching, cross-training, and competent handoff so a function can continue without one indispensable holder.

Validation Echo. Make feedback direct and usable. A social acknowledgment and an evidence-based review remain different records; supportive delivery does not establish correctness.

Simplification Push. Seek changes that reduce repeated effort. Include the cost of making, learning, checking, and maintaining the change. New knowledge need not simplify an unrelated process before it may be shared.

The tuning check

Ask what work is growing, what is deferred, which tool or skill is missing, whether agreed recovery time remains available, and whether a correction changed the next affected action. Voluntary personal reports can inform the review; no dashboard can infer how a person feels.

UBT records workload; OSIIN preserves knowledge; UPRR exposes resource state. Phase Lock v2.1 carries selected events between them. Allocation and governance remain separately declared.

Measure benefit before claiming renewal

A proposed comparison in hours

Define an equivalent task set, observation window, completion standard, and counting method before comparing a change. Let **B** be comparable baseline person-hours and **C** be current person-hours. Include coordination, learning, documentation, review, rework, and repair in the applicable totals.

Measure	Definition and limit
Workload difference	$B - C$, in person-hours. Keep estimates separate. Report changes in output, quality, scope, and deferred obligations.
Fractional difference	$(B - C) / B$, only when $B > 0$. This is a workload comparison, not a universal renewal coefficient.
Recovery opportunity	Hours made available under the work arrangement. Keep them separate from actual recovery and personal reports.
Service condition	Backlog, downtime, failures, coverage, missing data, and completion quality for the same declared interval.

Illustration only: if equivalent work previously took 12 person-hours and now takes 10 with the same counting basis, the difference is 2 person-hours. That does not establish that anyone took two hours of rest, that the change caused the difference, or that maintenance was not deferred.

The 1.2x objective and phase thresholds

The predecessor Phase Lock objective of 1.2 units of renewal per unit of knowledge integrity remains an **unvalidated historical target**. Its units, baseline, observation method, and attribution were not established. The older NTB threshold of +0.2 and CQ threshold near 0.85 are not adopted as qualification rules.

If Net Time Balance is used locally, define its inputs and units first. Do not double count overlapping task and rest intervals. Circulation Quality can be reported as separate latency, burden, coverage, access, and failure indicators; a composite score needs explicit construction and validation.

Evidence limit

A before-and-after comparison can identify a difference. Establishing why it occurred needs a suitable design and treatment of confounding factors. The historical phase names and 13th Harmonic language may remain cultural descriptions; they do not demonstrate a physical field or diagnose participant well-being.

Preserve what the work decision needs

The proposed **UBT-Work/2.1** profile links to a **Geometry-Claim/1** record when a workload, completion, or benefit claim becomes consequential. Phase Lock v2.1 defines the common profile; a short local note can implement it.

Geometry field	Application to UBT
Domain and retained view	Define eligible tasks and conditions D, and the record or summary pi used in the decision.
Required witnesses	Keep task completion, actual work, deferred work, service quality, upkeep, capacity and relevant admissibility as separately named outcomes W.
Evidence and criterion	Retain the original task, source/version, method, missing data, coverage, counterevidence and the declared acceptance rule.
Disposition and action	Separate relevant evidence r, explicit finding F, criterion-based disposition C, and performed action T. Attribute user adjudication separately from reviewer analysis.

Equal hours can lose the decisive difference

Constructed illustration: two work records both show eight person-hours and ten completed units. In one, scheduled maintenance was completed; in the other, it was deferred and equipment condition worsened. The hours/output summary cannot recover the upkeep witness. Retain maintenance state and the later capacity observation when the decision requires them. These are illustrative records, not a field result.

Let the evidence settle its own claim

For exact recovery, equal retained records must have equal required witnesses throughout the declared domain. A valid contrary pair establishes **FAILED / NOT_CLOSED** for that fixed claim. Clean partial coverage is **TESTED_ONLY**; insufficient evidence is **UNRESOLVED**; an unexecuted check is **NOT_RUN**. Complete valid coverage or a domain-wide proof can certify the declared claim, without certifying a whole work arrangement.

A failure can be established while repair remains pending. Carry a warranted disposition forward unless relevant evidence, scope, source validity, or the justified criterion changes. Consistency alone is insufficient: a constant but wrong label fails calibration. Acknowledgment and a revised plan remain distinct from the observed next action.

Output, upkeep, and the room to return

Shared context and present participation

Represent an encounter as **O and S interacting through their declared shared context L**, with time and dynamics recorded separately. L includes accessible records, language, tools and remembered corrections. It is a non-sovereign field, not another participant or an office that owns access. A chosen translator can work within it; field membership alone does not prove contact.

Retain four conditions separately: **contact, binding consent, usable exit, and absence of a compulsory third authority**. Unknown required conditions remain unresolved. More output or a positive resource balance cannot compensate for a failed condition. Prior participation or an inherited role does not manufacture present labor consent. Care does not create custody over a person; children are not a labor pool.

A resource budget with named units

For one resource, boundary and interval, use the proposed account:

Margin = available budget - task use - coordination use - maintenance use - regeneration use.

All entries use the same unit, and uses must be disjoint. This is a budget, not automatically a physical stock balance. Carry a remainder forward only for an actually storable resource with losses and transfers accounted for. Person-hours available in an interval do not become transferable stored time credits.

Capacity needs its own observation

For a named capacity, record its baseline, observed gains and losses, later condition, units or qualitative scale, interval, method and uncertainty. The identity **later capacity = initial capacity + gains - losses** organizes that account; it is not a predictive law. Spending time on regeneration does not establish that capacity recovered.

Keep useful output, maintenance burden, service quality and remaining capacity together. Extraction, maintenance, surplus and regeneration can describe different resources in the same interval; they are not a single automatic progression. A repeated output can conceal declining equipment condition, skill coverage or opportunities for recovery.

Return has more than one channel

The Geometry return record keeps changes in chart/observation, admissibility and retained context as a tuple. **MIXED** is valid. Equal observed output, full-state return and recovered function are different claims. Do not add incompatible units or infer an emotional state from the account. Physical energy and entropy also retain their own units and models.

Timing, context acquisition, and completion

Measure the event the task actually needs

Keep request onset, result readiness, release, first usable result and full completion as separate timestamps where observed. Record the clock, alignment and uncertainty; preserve raw event order, missing events and the applicable state version. An acknowledgment, spinner or transport receipt stops the task clock only if it meets the predeclared usable-outcome criterion.

Mean latency, standard deviation, duration count, interruption frequency, elapsed completion time and person-hours are different witnesses. The cyclic sequences (40,40,160,160) and (40,160,40,160), observed at their first entries, share mean 100, population SD 60 and current value 40, but have different next values. Their common summary is insufficient for that fixed-model next-step question.

Choose context under a declared acquisition policy

List observations actually available before the decision, their consent/access conditions, state version, acquisition time and fixed nonnegative costs where the finite model applies. Distinguish a fixed batch, a batch chosen after an already known label, and sequential queries. If no allowed observation separates a known incompatible pair, that menu is **UNREACHABLE** for the claim.

The v2.6 threshold construction needs $n-1$ batch queries versus an optimal worst-case $\text{ceil}(\log_2 n)$ sequential queries on its unchanged n -state domain. Every possible cut is still needed in the might-query inventory. Keep per-case query expense separate from maintaining tools, training and readiness. This exact mathematical advantage is not measured human-time savings. Live changing work or stock needs an explicit consistency or refresh rule.

Completion and presentation are different witnesses

Retain the original request, selected artifact and independently validated task result. Acquire the result from available pre-answer context. Preserving selected text under offline deletion does not prove live completion, less computation or a useful future response. Historical correction compliance stays attached to the original event.

No delay policy follows automatically

The Temporal v0.3 fixed-readiness model minimizes variance at a fixed feasible mean by delaying readiness R to $\max(R, T)$. Its assumptions include finite variance and causal release. Added waiting does not establish human benefit or lower total task time. Any live experiment must distinguish a delay/regularity tradeoff, an interaction, task quality and resource cost. No padding policy or new experiment is implemented here.

Use, verification, and revision record

A small operating cycle

Before: identify one bounded work set, existing backlog, required skills, voluntary capacity, resource dependencies, and a review interval. Keep the method light enough that measurement does not consume the benefit.

During: record scoped work and failures, offer acknowledgment, respect agreed limits, maintain coverage, and send corrections to the affected record or procedure.

At review: compare equivalent work, disclose missing data, inspect backlog and service quality, ask whether capacity and recovery opportunity improved, and choose one observed problem to address. No universal schedule, headcount, or score threshold is required.

Acceptance checks - NOT RUN in this edition

Demonstrate: no essential-access penalty for declining personal logs; duplicate events counted once; correction applied at the next eligible decision; unknown capacity preserved; workable succession, private/public separation and export; usable exit and emergency expiry. Also verify that upkeep survives an equal-hours summary, completed work meets the original criterion, and unrelated unknowns do not reopen a settled finding.

What this edition supersedes

The twelve named practices remain. This edition supersedes UBT v2.0 by integrating Geometry v2.6 into work records, claim disposition, shared context, capacity accounts, acquisition and timing. Automatic contribution rewards, compelled private renewal, and unvalidated field/benefit claims remain superseded. No historical result is relabeled as a new run.

Source basis and companions

Component source: **UBT (Universal Basic Time) quadrant tuning**, PDF copy dated 2026-09-13, three pages. Source SHA-256:

e52da8f77427aa6555675cb021f466c0ca1c1b9f3c01f1a560ba8ea516c77c16

Integration source: **The Abundance Engine - Open-Field Reference Architecture v7.0**, especially V.6, VII.2-VII.9, VIII, IX.7, and Appendix A. Source SHA-256:

6e620dc1b5feedbad0f182eece509c396667f8a37cf9a7bddd942d6a440f31d1

Related source: **UBTOSIINPhaseLockProtocol(2).md**. Companions: **OSIIN v3.1**, **Phase Lock v2.1**, **UPRR v3.1**. Source inventory, corrections and evidence limits: **Geometry_Integration_Review_2026-09-14.md**.

Geometry sources: **Working Master v2.6**; **Reflex Geometry v3**; **Temporal v0.3**; **Geometry + OMNIBUS 7**; and **Claim-specific closure geometry v0.1-r1**. OMNIBUS v7.79-r1 section 10 identifies its constitutional text as SUPPORTED RECONSTRUCTION.